

# Database Normalization and Joins

Charles Severance  
Barb Ericson



Python for Everybody  
[www.py4e.com/lectures3/](http://www.py4e.com/lectures3/)

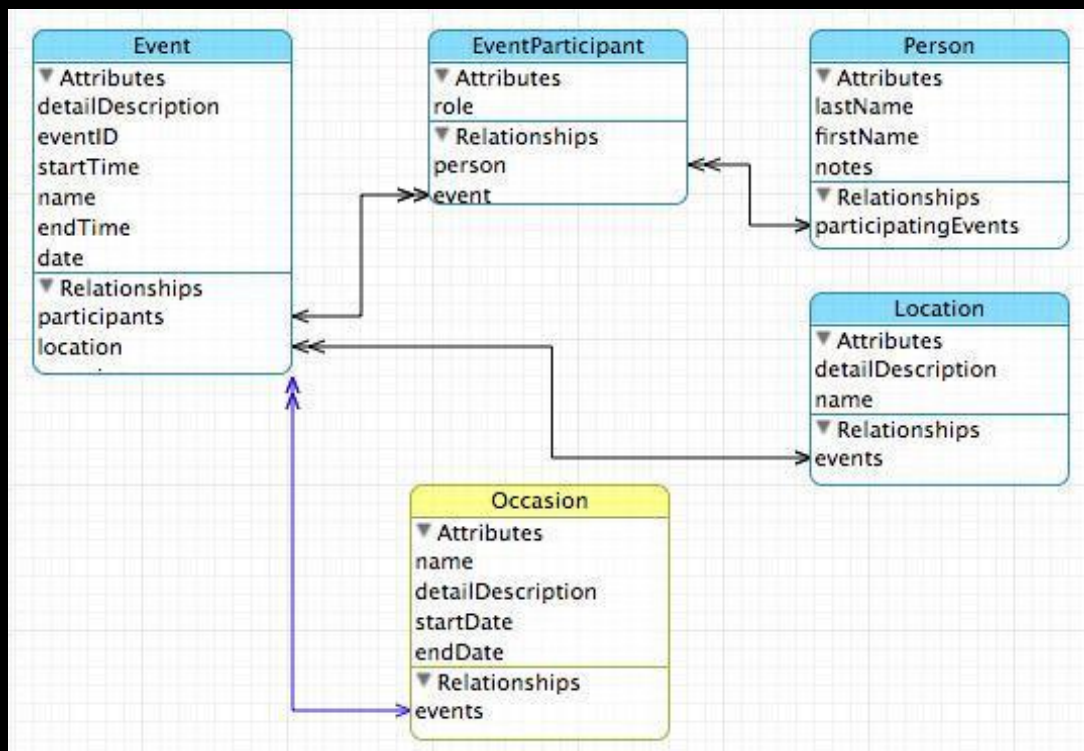


# Database Design

- Database design is an **art form** of its own with particular skills and experience
- Our goal is to avoid the really bad mistakes and design clean and easily understood databases
- Others may performance tune things later

# Building a Data Model

- Drawing a picture of the data objects for our application and then figuring out how to represent the objects and their relationships
- Basic Rule: Don't put the same string data in a database twice - use a relationship instead
- When there is one thing in the “real world” there should be one copy of that thing in the database



# Storing Movie Data

- What if we want to store data about movies in a database?
- We can use OMDb to get the data.

# What data do you need?

- Title, year, actors

Title	Year	Actors
Shaft	2000	"Samuel L. Jackson, Vanessa Williams, Christian Bale"
Glass	2019	"James McAvoy, Bruce Willis, Samuel L. Jackson"

# For each “piece of info”...

- Is the column an object or an attribute of another object?
- Once we define objects, we need to define the relationships between objects

Title                      Year

Actors

# Movie and Actor

- Break tables into the entities they represent

Movie

ID	Title	Year
1	Shaft	2000
2	Glass	2019

Actor

ID	Name
1	Samuel L. Jackson
2	Vanessa Williams
3	Christian Bale
4	James McAvoy
5	Bruce Willis



# Connecting Movie and Actor

- This is a many to many relationship
  - A movie can have many actors
  - An actor can appear in many movies

Movie_Actor	
movie_id	actor_id
1	1
1	2
1	3
2	4
2	5
2	1

# Music Tracks

- What if you want to store data for music tracks: title, length, artist, album, genre, rating, count?

Track	Len	Artist	Album	Genre	Rating	Count
<input checked="" type="checkbox"/> Hells Bells	5:13	AC/DC	Who Made Who	Rock	★★★★★	61
<input checked="" type="checkbox"/> Shake Your Foundations	3:54	AC/DC	Who Made Who	Rock	★★★★★	70
<input checked="" type="checkbox"/> Chase the Ace	3:03	AC/DC	Who Made Who	Rock		56
<input checked="" type="checkbox"/> For Those About To Rock (We ...	5:54	AC/DC	Who Made Who	Rock	★★★★★	61
<input checked="" type="checkbox"/> Guáimán	3:43	Altan	Natural Wonders M...	New Age		31
<input checked="" type="checkbox"/> Rode Across the Desert	4:10	America	Greatest Hits	Easy Listen...	★★★★★	23

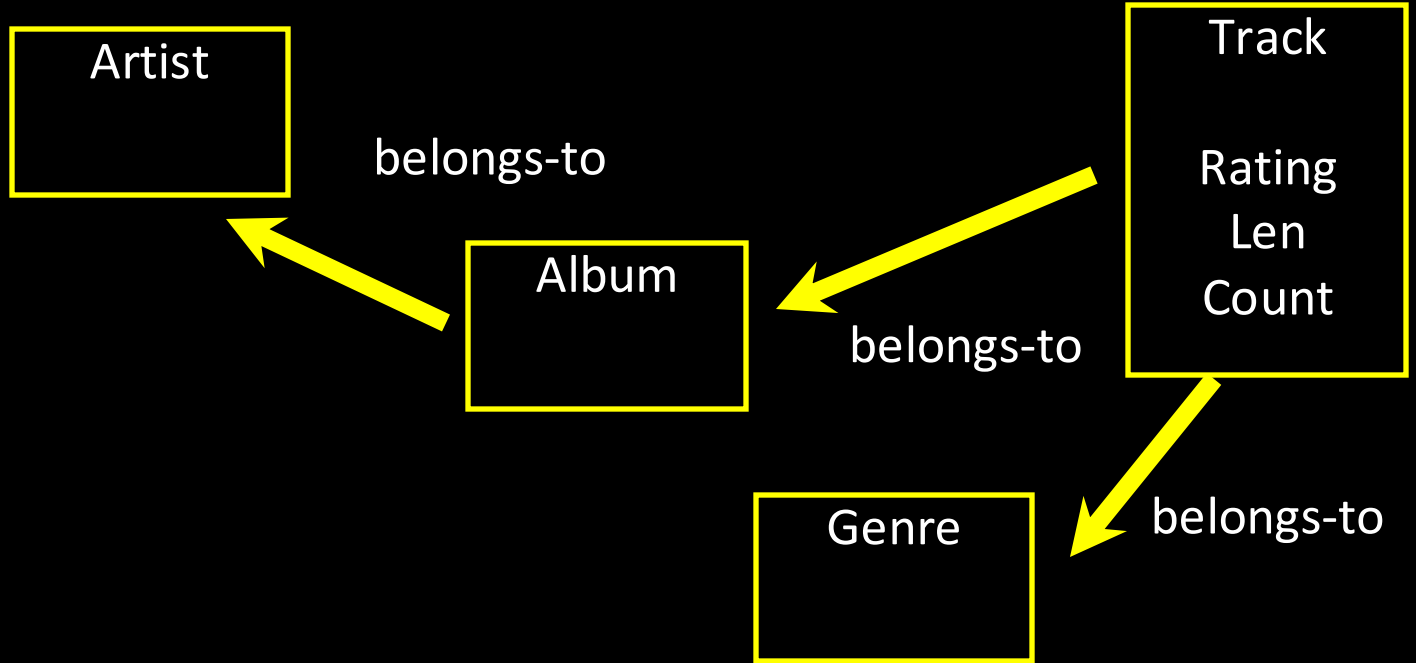
# For each “piece of info”...

- Is the column an object or an attribute of another object?
- Once we define objects, we need to define the relationships between objects

Len                      Album  
Genre  
Artist                      Rating  
Track                      Count

<input checked="" type="checkbox"/> Hells Bells	5:13	AC/DC	Who Made Who	Rock	★★★★★	61
<input checked="" type="checkbox"/> Shake Your Foundations	3:54	AC/DC	Who Made Who	Rock	★★★★★	70
<input checked="" type="checkbox"/> Chase the Ace	3:01	AC/DC	Who Made Who	Rock		56
<input checked="" type="checkbox"/> For Those About To Rock (We ...	5:54	AC/DC	Who Made Who	Rock	★★★★★	61
<input checked="" type="checkbox"/> Dúlamán	3:43	Altan	Natural Wonders M...	New Age		31
<input checked="" type="checkbox"/> Rode Across the Desert	4:10	America	Greatest Hits	Easy Listen...	★★★★★	23
<input checked="" type="checkbox"/> Now You Are Gone	3:08	America	Greatest Hits	Easy Listen...	★★★★★	18
<input checked="" type="checkbox"/> Tin Man	2:20	America	Greatest Hits	Easy Listen...	★★★★★	22

Track  
Album  
Artist  
Genre  
Rating  
Len  
Count



<input checked="" type="checkbox"/>	Hells Bells	5:13	AC/DC	Who Made Who	Rock	★★★★★	61
<input checked="" type="checkbox"/>	Shake Your Foundations	3:54	AC/DC	Who Made Who	Rock	★★★★★	70
<input checked="" type="checkbox"/>	Chase the Ace	3:01	AC/DC	Who Made Who	Rock		56
<input checked="" type="checkbox"/>	For Those About To Rock (We ...	5:54	AC/DC	Who Made Who	Rock	★★★★★	61
<input checked="" type="checkbox"/>	Dúlamán	3:43	Altan	Natural Wonders M...	New Age		31
<input checked="" type="checkbox"/>	Rode Across the Desert	4:10	America	Greatest Hits	Easy Listen...	★★★★★	23
<input checked="" type="checkbox"/>	Now You Are Gone	3:08	America	Greatest Hits	Easy Listen...	★★★★★	18
<input checked="" type="checkbox"/>	Tin Man	2:20	America	Greatest Hits	Easy Listen...	★★★★★	22

# Representing Relationships in a Database

# Database Normalization (3NF)

- There is \*tons\* of database theory - way too much to understand without excessive predicate calculus
- Do not replicate string data - reference it instead
- Use integers for keys and for references
- Add a special “key” column to each table which we will make references to. By convention, many programmers call this column “id”

[http://en.wikipedia.org/wiki/Database\\_normalization](http://en.wikipedia.org/wiki/Database_normalization)

<input checked="" type="checkbox"/> Hells Bells	5:13	AC/DC	Who Made Who	Rock	★★★★★	61
<input checked="" type="checkbox"/> Shake Your Foundations	3:54	AC/DC	Who Made Who	Rock	★★★★★	70
<input checked="" type="checkbox"/> Chase the Ace	3:01	AC/DC	Who Made Who	Rock		56
<input checked="" type="checkbox"/> For Those About To Rock (We ...	5:54	AC/DC	Who Made Who	Rock	★★★★★	61
<input checked="" type="checkbox"/> Dúlamán	3:43	Altan	Natural Wonders M...	New Age		31
<input checked="" type="checkbox"/> Rode Across the Desert	4:10	America	Greatest Hits	Easy Listen...	★★★★★	23
<input checked="" type="checkbox"/> Now You Are Gone	3:08	America	Greatest Hits	Easy Listen...	★★★★★	18
<input checked="" type="checkbox"/> Tie Me	3:30	America	Greatest Hits	Easy Listen...	★★★★★	22

We want to keep track of which band is the “creator” of each music track...

What album does this song “belong to”??

Which album is this song related to?

# Integer Reference Pattern


We use integers to  
reference rows in  
another table

id	name
Filter	Filter
1	Led Zepplin
2	AC/DC

Artist

id	artist_id	title
Filter	Filter	Filter
1	2	Who Made Who
2	1	IV


Album





# Three Kinds of Keys

- **Primary key** - generally an integer auto-incremented field
- **Logical key** - What the outside world uses for lookup
- **Foreign key** - generally an integer key pointing to a row in another table



# Key Rules

## Best practices

- Never use your **logical key** as the **primary key**
- **Logical keys** can and do change, albeit slowly
- **Relationships** that are based on matching string fields are less efficient than integers

```
User
id
login
password
name
email
created_at
modified_at
login_at
```

# PI #1

Q-1: Which of the following is a foreign key in the `tracks` table given that the table has a `TrackId` of 1, a `Name` of "Dog Eat Dog" a `Length` of 2.30 and a `GenreId` of 1

- ☐ A. Name
- ☐ B. TrackId
- ☐ C. Length
- ☐ D. GenreId

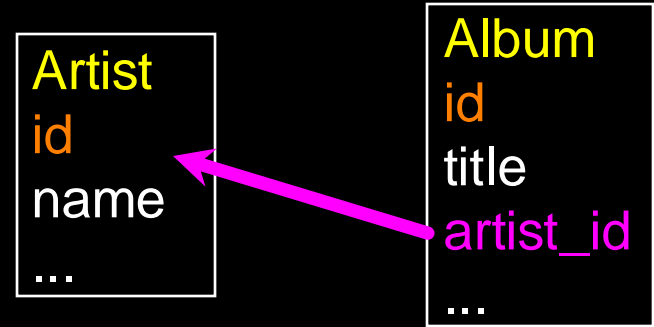
Check Me

Compare me

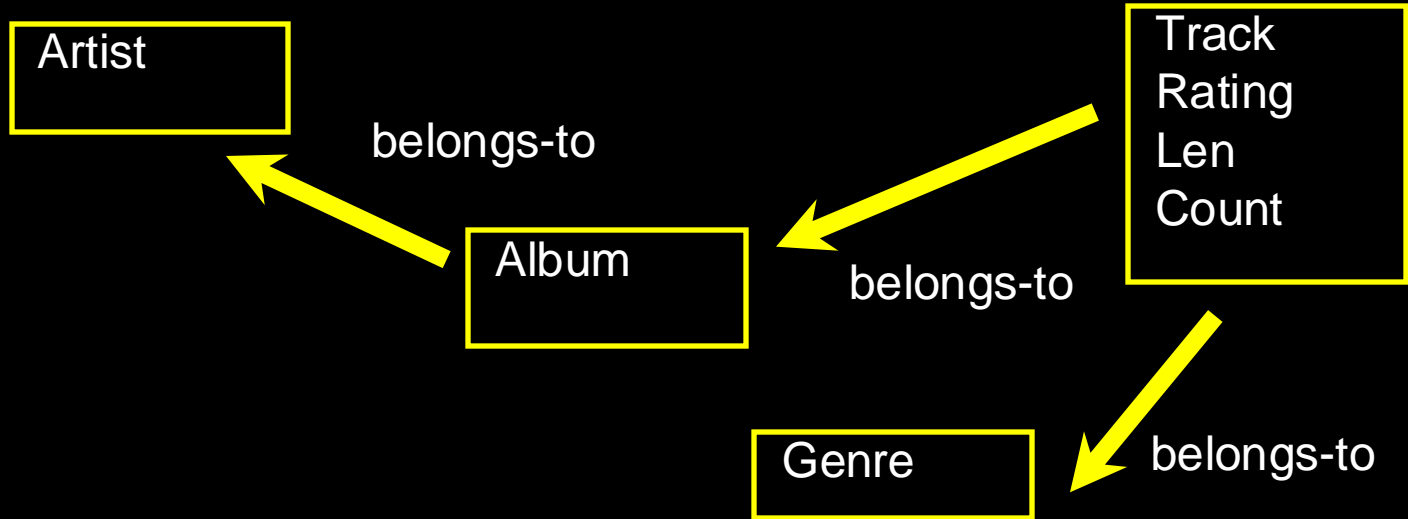
Problem: 1 -- Activity: 1 Multiple Choice (pi-db-join-foreign-key-tracks)

# Foreign Keys

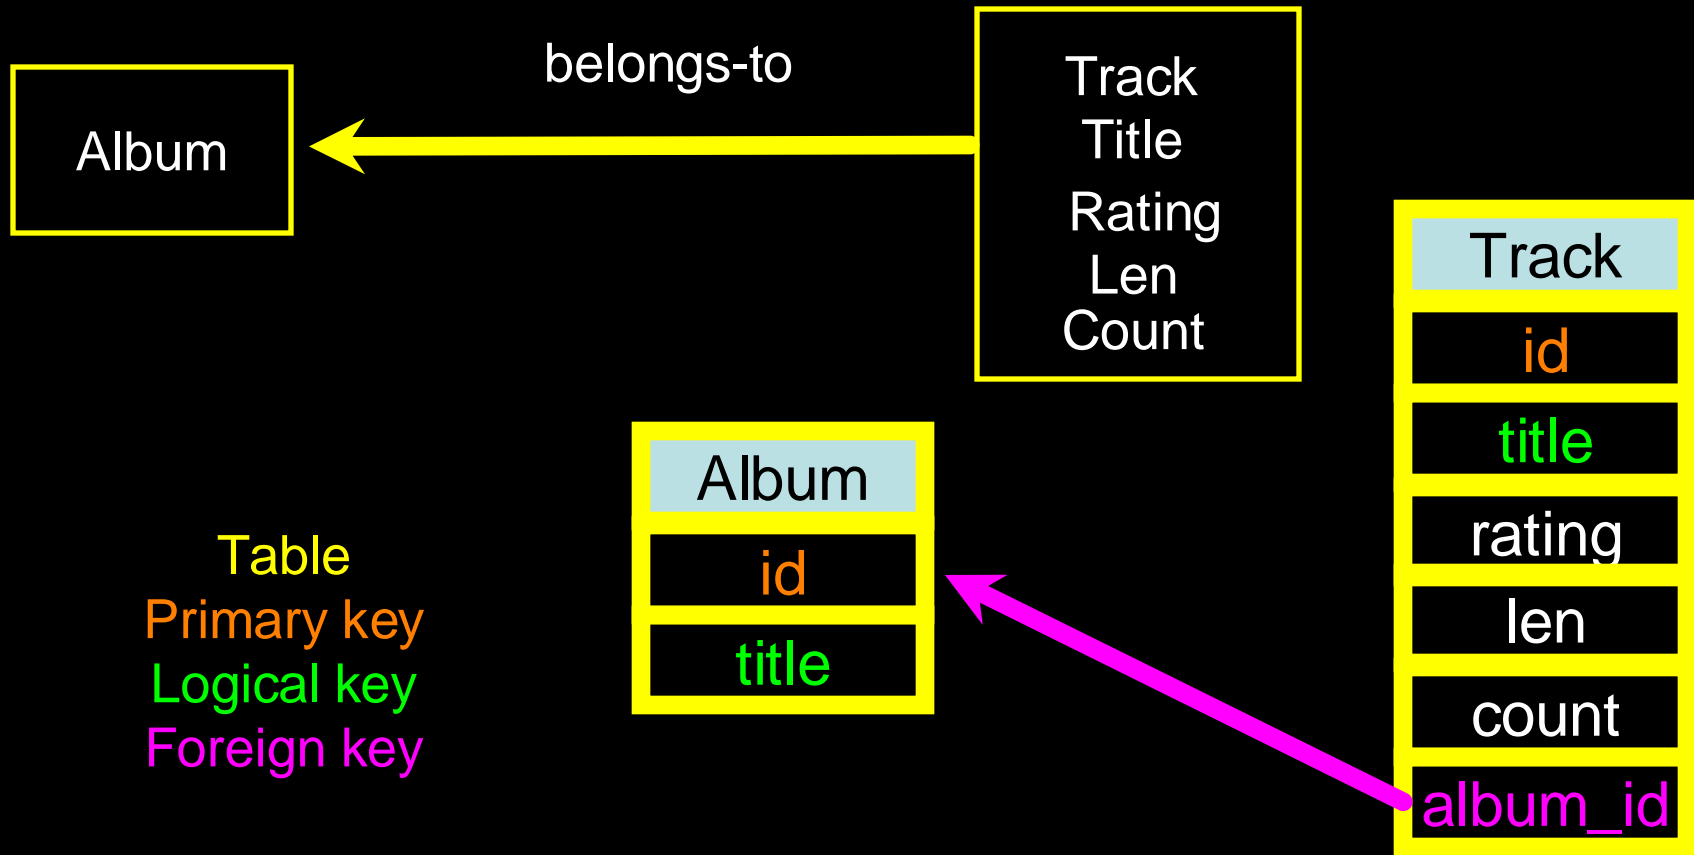
- A **foreign key** is when a table has a column that contains a key which points to the **primary key** of another table.
- When all primary keys are integers, then all foreign keys are integers - this is good - very good



# Relationship Building (in tables)



<input checked="" type="checkbox"/> Hells Bells	5:13	AC/DC	Who Made Who	Rock	★★★★★	61
<input checked="" type="checkbox"/> Shake Your Foundations	3:54	AC/DC	Who Made Who	Rock	★★★★★	70
<input checked="" type="checkbox"/> Chase the Ace	3:01	AC/DC	Who Made Who	Rock		56
<input checked="" type="checkbox"/> For Those About To Rock (We ...	5:54	AC/DC	Who Made Who	Rock	★★★★★	61
<input checked="" type="checkbox"/> Dúlamán	3:43	Altan	Natural Wonders M...	New Age		31
<input checked="" type="checkbox"/> Rode Across the Desert	4:10	America	Greatest Hits	Easy Listen...	★★★★★	23
<input checked="" type="checkbox"/> Now You Are Gone	3:08	America	Greatest Hits	Easy Listen...	★★★★★	18
<input checked="" type="checkbox"/> Tin Man	3:20	America	Greatest Hits	Easy Listen...	★★★★★	22



Artist
id
name




Album
id
title
artist_id

Track
id
title
rating
len
count
album_id
genre_id

Genre
id
name

Table  
Primary key  
Logical key  
Foreign key

Naming FK artist\_id is a  
convention





Edit table definition

Table

Artist

Advanced

Fields Constraints

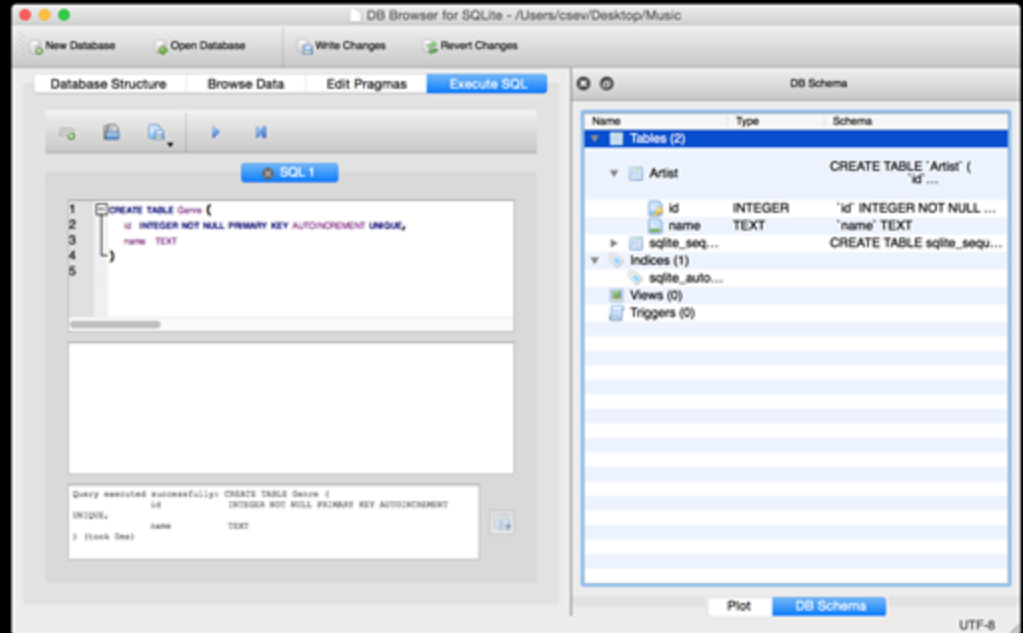
Add Remove Move to top Move up Move down Move to bottom

Name	Type	NN	PK	AI	U	Default	Check
id	INTEGER	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>		<input checked="" type="checkbox"/>

```
1 CREATE TABLE "Artist" (  
2   "id" INTEGER NOT NULL UNIQUE,  
3   PRIMARY KEY("id" AUTOINCREMENT)  
4 );
```

Cancel OK

NN = Not null  
PK = Primary Key  
AI = Autoincrement  
U = Unique



```
CREATE TABLE Genre (  
    id        INTEGER NOT NULL PRIMARY KEY AUTOINCREMENT UNIQUE,  
    name      TEXT  
)
```

```
CREATE TABLE Album (  
    id          INTEGER NOT NULL PRIMARY KEY AUTOINCREMENT UNIQUE,  
    artist_id   INTEGER,  
    title       TEXT  
)
```

```
CREATE TABLE Track (  
    id          INTEGER NOT NULL PRIMARY KEY AUTOINCREMENT UNIQUE,  
    title       TEXT,  
    album_id    INTEGER,  
    genre_id    INTEGER,  
    len         INTEGER, rating INTEGER, count INTEGER  
)
```

id	title	album_id	genre_id	len	rating	count
Filter	Filter	Filter	Filter	Filter	Filter	Filter
1	Black Dog	2	1	297	5	0
2	Stairway	2	1	482	5	0
3	About to Rock	1	2	313	5	0
4	Who Made Who	1	2	207	5	0

Track

Album

id	artist_id	title
Filter	Filter	Filter
1	2	Who Made Who
2	1	IV

Artist

id	name
Filter	Filter
1	Led Zepplin
2	AC/DC

id	name
Filter	Filter
1	Rock
2	Metal

Genre

# Peer Instruction #2

Q-1: How many database tables would best represent the data for a restaurant such as Name: "Spencer", Cost: \$\$, Stars: 4.5, Num Reviews: 168, Food Type: American (New)

- ☐ A. 1
- ☐ B. 2
- ☐ C. 3
- ☐ D. 4

Check Me

Compare me

Problem: 2 -- Activity: 1 Multiple Choice (pi-db-join-how-many-tables)

# Using Join Across Tables

[http://en.wikipedia.org/wiki/Join\\_\(SQL\)](http://en.wikipedia.org/wiki/Join_(SQL))

# Relational Power

- By removing the replicated data and replacing it with references to a single copy of each bit of data we build a “web” of information that the relational database can read through very quickly - even for very large amounts of data
- Often when you want some data it comes from a number of tables linked by these foreign keys

# The JOIN Operation

- The JOIN operation **links across several tables** as part of a select operation
- You must tell the JOIN **how to use the keys** that make the connection between the tables using an **ON clause**




# Types of Joins

## Different Types of SQL JOINS

Here are the different types of the JOINS in SQL:

- **(INNER) JOIN** : Returns records that have matching values in both tables
- **LEFT (OUTER) JOIN** : Returns all records from the left table, and the matched records from the right table
- **RIGHT (OUTER) JOIN** : Returns all records from the right table, and the matched records from the left table
- **FULL (OUTER) JOIN** : Returns all records when there is a match in either left or right table



id	artist_id	title
Filter	Filter	Filter
1	2	Who Made Who
2	1	IV

Album

Artist

	title	name
1	Who Made Who	AC/DC
2	IV	Led Zeppelin

id	name
Filter	Filter
1	Led Zeppelin
2	AC/DC

`SELECT Album.title, Artist.name FROM Album JOIN Artist ON Album.artist_id = Artist.id`

What we want  
to see

The tables that  
hold the data

How the tables  
are linked

Album

id	artist_id	title
Filter	Filter	Filter
1	2	Who Made Who
2	1	IV

Artist

id	name
Filter	Filter
1	Led Zeppelin
2	AC/DC

Result

	title	artist_id	id	name
1	Who Made Who	2	2	AC/DC
2	IV	1	1	Led Zeppelin

SELECT Album.title, Album.artist\_id, Artist.id,  
 Artist.name  
 FROM Album JOIN Artist ON Album.artist\_id = Artist.id

Result

	title	name
1	Black Dog	Rock
2	Stairway	Rock
3	About to Rock	Metal
4	Who Made Who	Metal

Track

id	title	album_id	genre_id	len	rating	count
Filter	Filter	Filter	Filter	Filter	Filter	Filter
1	Black Dog	2	1	297	5	0
2	Stairway	2	1	482	5	0
3	About to Rock	1	2	313	5	0
4	Who Made Who	1	2	207	5	0

Genre

id	name
Filter	Filter
1	Rock
2	Metal

SELECT Track.title, Genre.name FROM Track JOIN Genre ON Track.genre\_id = Genre.id

What we want  
to see

The tables that  
hold the data

How the tables  
are linked

SELECT Track.title, Artist.name, Album.title,  
Genre.name FROM Track JOIN Genre JOIN Album  
JOIN Artist ON Track.genre\_id = Genre.id and  
Track.album\_id = Album.id and Album.artist\_id =  
Artist.id

	title	name	title	name
1	Black Dog	Led Zeppelin	IV	Rock
2	Stairway	Led Zeppelin	IV	Rock
3	About to Rock	AC/DC	Who Made Who	Metal
4	Who Made Who	AC/DC	Who Made Who	Metal

What we want to see

The tables which hold  
the data

How the tables are  
linked

# Summary

- Relational databases allow us to **scale** to very large amounts of data
- The key is to have **one copy of any string data** element and use relations and joins to link the data to multiple places
- This greatly **reduces the amount of data which much be scanned** when doing complex operations across large amounts of data
- Database and SQL design is a bit of an **art form**

# Lecture Participation

- Go to the ebook and do Practice-Join for up to 10 points



## Acknowledgements / Contributions



These slides are Copyright 2010- Charles R. Severance ([www.dr-chuck.com](http://www.dr-chuck.com)) of the University of Michigan School of Information and [open.umich.edu](http://open.umich.edu) and made available under a Creative Commons Attribution 4.0 License. Please maintain this last slide in all copies of the document to comply with the attribution requirements of the license. If you make a change, feel free to add your name and organization to the list of contributors on this page as you republish the materials.

...

Initial Development: Charles Severance, University of Michigan School of Information

... Insert new Contributors here